Gmail

AI Fusion Labs <aifusionlabs@gmail.com>

Thanks for joining Anam!

6 messages

Ben Carr <ben@hq.anam.ai> Tue, Feb 17, 2026 at 9:55 AM
Reply-To: info@anam.ai
To: aifusionlabs@gmail.com

Welcome to Anam

Now you have joined, lets get you started!

Depending on what you had in mind, you might find the best place to start being by grabbing yourself an API Key.

Just click the link below to get straight there!

Get your API Key

Already have your Key?

Lets get building, the true value of Anam comes when you connect it to your application.

Install the SDK to get started - just a few lines of code and you will be on your way!

See the docs

Need some help?

if you need any assistance at all, feel free to reply or post on our Slack Community!

Happy creating!

Join us on Slack

Book a Demo

Anam.ai LTD, White Collar Factory, Old Street Yard, London, London
EC1Y 8AF, United Kingdom, 07384669569

Unsubscribe Manage preferences

AI Fusion Labs <aifusionlabs@gmail.com> Tue, Feb 17, 2026 at 11:22 AM
To: info@anam.ai

Hello Ben<
I'm Rob with AI Fusion Labs, and we are building and demoing a real-time video avatar + voice agent platform across multiple verticals (field service, legal, customer support, etc.).

Today we're using **Tavus CVI** and the results are strong, but I'm intrigued by what **Anam** can offer and I'd love to evaluate it seriously.

A few questions I'm hoping you can help with:

1. **Video quality/resolution**

- Your API page mentions **25fps at 720×480**. Do you have a higher-res mode or roadmap (e.g., 720p/1080p), and what's the trade-off in terms of latency/cost?

2. **Costing + bootstrapped developer reality**

- During testing and iterative demos, we burn minutes fast. Are there **developer/testing credits**, a startup program, or a way to get a larger minute pool for evaluation (especially while we're still bootstrapping)?

3. **Scale/concurrency**

- Your pricing page shows simultaneous sessions scaling by tier (e.g., 1/3/5/30+). How do you recommend planning concurrency for production usage, and what do typical enterprise concurrency limits look like?

4. **Integration effort (tools/webhooks/backend)**

- We can handle integration work, but we need to scope it correctly. Since Anam supports tool calling (client tools, knowledge tools, and webhook tools), what's the "minimum viable" backend setup you see for a production pilot?

- If you have a reference architecture (recommended event handling patterns, webhook security, retries, observability), I'd love to review it.

Thanks — excited to try Anam and see if it can become our scalable/efficient avatar layer.

Best,
Rob
[Quoted text hidden]

Caoimhe Murphy <caoimhe@anam.ai> Tue, Feb 17, 2026 at 1:55 PM
To: AI Fusion Labs <aifusionlabs@gmail.com>

Hey - thanks for the email. Here are some replied to your questions below.

A few questions I'm hoping you can help with:

1. **Video quality/resolution**

- Your API page mentions **25fps at 720×480**. Do you have a higher-res mode or roadmap (e.g., 720p/1080p), and what's the trade-off in terms of latency/cost? Yes our next model, the resolution will improve which will launch in Q2 this year.

2. **Costing + bootstrapped developer reality**

- During testing and iterative demos, we burn minutes fast. Are there **developer/testing credits**, a startup program, or a way to get a larger minute pool for evaluation (especially while we're still bootstrapping)? We have a start up plan which is 1K per month for 1000 minutes and full enterprise access, as a start up ourselves we dont yet offer free credits, we also have generous no of minutes across our other self service plans that might work for your use case

3. **Scale/concurrency**

- Your pricing page shows simultaneous sessions scaling by tier (e.g., 1/3/5/30+). How do you recommend planning concurrency for production usage, and what do typical enterprise concurrency limits look like? its very use case dependant eg a website video agent vs a training agent that has 1000 people using at 3pm each day if you share more about your use case I can give you a better ratio

4. **Integration effort (tools/webhooks/backend)**

- We can handle integration work, but we need to scope it correctly. Since Anam supports tool calling (client tools, knowledge tools, and webhook tools), what's the "minimum viable" backend setup you see for a production pilot? You can use the lab today without any deep integration effort to have something up and running for a production pilot

- If you have a reference architecture (recommended event handling patterns, webhook security, retries, observability), I'd love to review it. let me check and follow up

Thanks,

C

[Quoted text hidden]

Rudy Wilson-Evans <rudy@anam.ai> Wed, Feb 18, 2026 at 7:47 AM
To: AI Fusion Labs <aifusionlabs@gmail.com>

Hey there Rob,

Currently we dont offer additional resolutions, though this on our roadmap alongside the updates for improved emotional control and active listening.
We can support you to get started as we have an offer of a month free with our latest model launch with the code CARA-3 which would allow you to use any plan for free for a full month, if you are looking for more minutes beyond that it would be great to understand your needs in more details.

We currently support customers with 100s of concurrent sessions but do not have a top limit today we would make sure

though that we add you to an appropriate enterprise tier to account for your needs and then assess and scale as needed for you.

Depending on the way you are looking to handle the tool calling we can share some best practices but many clients have hooked into their own LLM for that however you are able to handle this on the ANam side as well as you need.

It could be worth a chat with our team if you are interested in learning more, otherwise i can share over our cookbooks and see if that helps at all

Thanks

Rudy

[Quoted text hidden]

AI Fusion Labs <aifusionlabs@gmail.com> Thu, Feb 26, 2026 at 5:26 PM
To: Rudy Wilson-Evans <rudy@anam.ai>

Rudy — thanks, sorry for the delay in getting back to you. We are actively testing and seeing strong results, and some occasional blips also.
One major concern I have is that the voice has a tendency to change, so much so it is recognizable. I was able to capture an instance in this video, link below, and it happened at only the beginning, which seems to be the only trend I can share. I have tried to change voice models several times, but I still keep running into it. I assume it is on my end, possibly in my system prompt (maybe system bloat, not sure what your context window size is), but no luck after several revisions.
https://www.loom.com/share/23dd2614c84440b38751d453a24ab977
The site if you want to visit and test. I may be changing things along the way, though....
https://x-agent-website-a.vercel.app/

I cannot tell you how much the offer of a month free is to us. Unlimited testing is huge, as currently we are sinking the majority of our 'minutes' in testing alone. Some of the other points you mentioned are mainly for my roadmap and some current prospects I am working with. For the time being, we'll keep focusing on testing, and if your offer for your 'cookbooks' still stands, that would be incredible! Always looking to learn and see new ways to approach utilizing Anam.

Happy to jump on a quick call, even just to connect and talk 'shop'.

Best,
Rob

[Quoted text hidden]

Rudy Wilson-Evans <rudy@anam.ai> Fri, Feb 27, 2026 at 3:59 AM
To: AI Fusion Labs <aifusionlabs@gmail.com>
Cc: Anam <support@anam.ai>

Hey Rob,

Thanks for the update and for sharing the video and site link - super helpful.

We do see some variability with ElevenLabs voices from time to time but you can use https://docs.anam.ai/concepts/personas#elevenlabs-voice-configuration to configure the stability but if you need additional support please let us know.

Glad the free month is useful - let me know if you need it extended or if you hit any other limits.

Our Cookbooks are all here for you: https://anam.ai/cookbook

Appreciate you testing so thoroughly and always keen to see what you're building.

Rudy

[Quoted text hidden]